(E0 261) Database Management Systems
MINIREL Project Description

1 OVERVIEW

In this project, you will implement a simplified single-user relational database system, called **MINIREL**. The MINIREL project involves writing code for both the logical layer and the physical layer of the dbms. For ease of implementation, the logical layer is split into two sub-layers, an algebra layer and a schema layer. Therefore, the MINIREL architecture is a three-layer hierarchy comprised of an **Algebra Layer**, a **Schema Layer** and a **Physical Layer**. Your implementation will have to reflect this hierarchy.

The Algebra Layer implements the main relational algebra constructs such as *select*, *project*, etc. The Schema Layer implements routines that define and manipulate the database schema. Utility routines for manipulating relations are also included in this layer. The Physical Layer implements routines which directly manipulate files and records through the operating system interface. The meta-data information is stored in **catalogs**.

You will implement MINIREL *bottom-up*, that is, starting with the Physical Layer, then writing the Schema Layer, and finally, the Algebra Layer. We will provide a **front-end** to translate user dbms commands into calls to MINIREL. This front-end will parse the dbms commands and call routines of your Algebra Layer and Schema Layer to process each command. During program development, you should use the front-end for writing queries to test the correctness and robustness of your programs.

1.1 Programming Environment

MINIREL is to be implemented in C or C++ on the Unix operating system. You can develop the program on any machine you like but your final demo will have to execute on a Linux box.

1.2 Project Submission

To help pace yourself on the MINIREL project, the following 12-week schedule is recommended:

Project Design:	one week
Physical Layer:	three weeks
Schema Layer:	two weeks
Algebra Layer:	two weeks
Index Routines:	three weeks
Testing:	one week

You are expected to hand in a fully commented listing for each layer of your program. A documentation template and a sample documentation is available in the course account. Similarly, sample dbms command files are available for testing the correctness and robustness of your programs. After project completion, each team will have to provide an online demo.

1.3 Grading

For this project, only half the credit will be for program correctness — the remaining is for documentation, style, efficiency, and robustness, so you should write *quality* code. Typically, the team grade will simply be assigned to each team member. However, in cases where the work distribution is noticeably skewed, we will adjust your individual grades accordingly. Since questions on the project may appear in the exams and during the demo, make sure that you are familiar with the entire code. Copying of code across teams will instantly result in *all* concerned failing the course, so it is your responsibility to prevent others from accessing your code.

2 THE FRONT-END

This section describes the front-end system (which we will provide in both C and C++ for a variety of architectures) and the set of commands that can be input by users to the front-end. The front-end consists of a *parser* which parses user commands and converts them to calls to your logical layer (Algebra Layer + Schema Layer) routines. The parser uses the standard **(argc, argv)** technique to pass parameters to your procedures. All parameters passed into your routines are in ASCII (character string) form, so they have to be converted to the appropriate types before processing. You *must* stick to the procedure call interface specified for the logical layer routines in Sections 3 and 4 in order to be able to use the front-end system.

The front-end verifies the *syntactic* correctness of user requests. However, the front-end has no knowledge about what relations exist, what attributes the relations contain, or the types of their attributes. Therefore, your routines must detect and deal with the possibility that incoming requests may make incorrect references to relations and attributes. It is up to your routines to do *all* of the semantic error detection and handling. Ideally, you should detect all semantic errors before you make any changes to the database, as undoing such changes will be harder than preventing them in the first place. Also, you should enforce meta-data security, that is, not allow users to directly alter the catalogs. Finally, your routines should return an error code and print out *meaningful* error messages when errors do occur.

2.1   USER COMMANDS

The commands that MINIREL users can input to the front-end are listed below. The syntax of the commands is explained in the notes following the command list.

1. `createdb DBNAME;`

2. `destroydb DBNAME;`

3. `opendb DBNAME;`

4. `closedb;`

5. `quit;`

6. `create RELATION_NAME ( ATTR_NAME = FORMAT [ , ATTR_NAME = FORMAT ]* );`

7. `destroy RELATION_NAME;`

8. `load RELATION_NAME from FILENAME;`

9. `print RELATION_NAME;`

10. `buildindex for RELATION_NAME on ATTR_NAME;`

11. `dropindex for RELATION_NAME [ on ATTR_NAME ];`

12. `select into RELATION_NAME from RELATION_NAME where ( ATTR_NAME OP VALUE );`

13. `project into RELATION_NAME from RELATION_NAME ( ATTR_NAME [ , ATTR_NAME ]* );`

14. `join into RELATION_NAME ( RELATION_NAME . ATTR_NAME, RELATION_NAME . ATTR_NAME );`

15. `insert into RELATION_NAME ( ATTR_NAME = VALUE [ , ATTR_NAME = VALUE ]* );`

16. `delete from RELATION_NAME where ( ATTR_NAME OP VALUE );`

NOTES

1. All lower case words are key words.

2. The upper case words have the following interpretation:

RELATION_NAME	–	string
ATTR_NAME	–	string
DBNAME	–	string
FORMAT	–	string
FILENAME	–	string
VALUE	–	quoted string or number
OP	–	comparison operators like $<$, $>=$ etc.

where

string: a character set that starts with an alphabet and is followed by an
arbitrary number of alphabets and digits

quoted string: a string within double quotes (semicolon is not allowed within the quotes)

number: integer or real number

operators: $<, <=, <>, >, >=, =$

3. All the strings and numbers passed to the routines are in ASCII **except** the comparison operators which are integer constants. They are as follows:

OP	CONSTANT
=	501
>=	502
>	503
<=	504
<>	505
<	506

4. Terms within square brackets '[ ]' are optional. A '*' after the right square bracket ']' means terms within the square brackets '[ ]' can be repeated any number of times.

5. All commands to the front-end must end with a ';'.

3   ALGEBRA LAYER

This section of the handout describes the routines for the Algebra Layer of MINIREL, and how they are invoked by the front-end. Each of the Algebra Layer routines implements a basic relational algebra command. The commands are *select, project, join, insert,* and *delete.* Each of these routines takes one or two relations as the input and yields a new relation as its result. Result relations are included in the database.

1. *Select(argc, argv)*:

argv[0]	= "select"
argv[1]	= result relation
argv[2]	= source relation
argv[3]	= attribute name
argv[4]	= operator
argv[5]	= value
argv[argc]	= NIL

   This routine implements the relational selection command. It first creates a result relation with the same set of attributes as the source relation. It then applies the selection criterion to the records in the source relation, and places the selected records in the result relation. In MINIREL, only one selection criterion is handled at a time. The selection criterion itself is specified via the name of an attribute, a comparison operator, and a value. The supported operators are $<, >, <=, >=, =,$ and $<>$; as described in Section 2, the operators are passed to MINIREL by the front-end as integers.

2. *Project (argc, argv)*:

argv[0]	= "project"
argv[1]	= result relation
argv[2]	= source relation
argv[3]	= attribute name 1
argv[4]	= attribute name 2
. . .	. . .
argv[argc-1]	= attribute name N
argv[argc]	= NIL

   This routine implements the relational projection command. It first creates the result relation with the attributes given in *argv*, then performs the projection. Attributes are matched up by name, not by position, and the order of attributes in the result relation is the order specified in *argv*. Since eliminating attributes may introduce duplicate tuples, your implementation must ensure that the duplicates are removed from the output relation.

3. *Join(argc, argv)*:

argv[0]	= "join"
argv[1]	= result relation
argv[2]	= source relation 1
argv[3]	= attribute name 1
argv[4]	= source relation 2
argv[5]	= attribute name 2
argv[argc]	= NIL

This routine implements the relational join command. For simplicity, MINIREL only implements *natural join,* not the full join command. This routine first creates the specified result relation, then performs a join on the two source relations based on either the **nested loops** technique or the **index join** technique.

Since natural join is implemented, the join attribute appears only *once* in the result relation. The join attribute from the first source relation is used as the result's join attribute. The result relation's attribute order is: *source 1 attributes* (including the join attribute) followed by *source 2 attributes*. While joins on unlike field types (e.g., an integer field and a string field) are errors, strings of different lengths are accepted as being of the same type for the purpose of performing joins. Finally, you must deal with the possibility that the two source relations may have one or more attributes with the same name — devise an attribute renaming scheme to resolve such conflicts. Note, however, that the join attribute should *not* be renamed.

4. *Insert(argc, argv)*:

argv[0]	= "insert"
argv[1]	= relation name
argv[2]	= attribute name 1
argv[3]	= attribute value 1
. . .	. . .
argv[argc-2]	= attribute name N
argv[argc-1]	= attribute value N
argv[argc]	= NIL

This routine implements the relational insert command, adding the tuple given in *argv* to the named relation. Since the input values are in ASCII, the appropriate conversion is performed before entering into the database. Attributes are matched up by name, not by their order. The length of incoming string-type data is regulated by padding with NULLs if too short, or truncating if too long. Make sure that adding the tuple does not introduce a duplicate. If the attribute set is incorrectly specified, an error message should be generated.

5. *Delete(argc, argv)*:

argv[0]	= "delete"
argv[1]	= source relation
argv[2]	= attribute name
argv[3]	= operator
argv[4]	= value
argv[argc]	= NIL

This routine implements the relational delete command. It is similar to *Select*, but it deletes the specified records from the source relation instead of placing them in a result relation. You should be able to implement this routine quickly by copying your *Select* code and then making a few changes to it.

Note: If appropriate indexes exist, MINIREL should use them to process the above operations efficiently.

4 SCHEMA LAYER

In the Schema Layer, MINIREL implements routines that form part of the logical layer and deal with schema definition and manipulation. Utility routines for printing the contents of relations, loading relations from binary data files, and building indexes on relations, are also included in this layer. This section describes each of the routines in the Schema Layer. Each routine is called by the front-end or by Algebra Layer routines with *(argc, argv)* parameters (as in the Algebra Layer).

1. *CreateDB(argc, argv)*:

```
argv[0]    = "createdb"
argv[1]    = database name
argv[argc] = NIL
```

This routine creates a new database. It creates the relation and attribute catalogs and loads them with the appropriate initial information. You should check that the given database directory does not already exist.

2. *DestroyDB(argc, argv)*:

```
argv[0]    = "destroydb"
argv[1]    = database name
argv[argc] = NIL
```

This routine first checks that the database has been closed. If it is not closed, it closes it and then destroys the database.

Note: Only the files and directories related to the specified database should be destroyed, without affecting other databases.

3. *OpenDB(argc, argv)*:

```
argv[0]    = "opendb"
argv[1]    = database name
argv[argc] = NIL
```

This routine opens a database for subsequent use. It changes the working directory to be the database directory, opens the catalogs and initializes the various global data structures that are used by other MINIREL routines.

4. *CloseDB(argc, argv)*:

```
argv[0]    = "closedb"
argv[argc] = NIL
```

This routine closes the currently opened database and changes the working directory to the original directory from which MINIREL was invoked.

5. *Quit()*:

```
argv[0]    = "quit"
argv[argc] = NIL
```

This routine first checks that the database has been closed. If it is not closed, it closes it and then terminates the programs.

Note: For CreateDB, DestroyDB, and OpenDB, assume that the full path name is given in the database name parameter. At the end of the query interaction, your program should <u>return</u> to the directory from which you invoked MINIREL.

6. *Create(argc, argv)*:

```
argv[0]      = "create"
argv[1]      = relation name
argv[2]      = first attribute name
argv[3]      = first attribute format
 . . .          . . .
argv[argc-2] = Nth attribute name
argv[argc-1] = Nth attribute format
argv[argc]   = NIL
```

This routine creates a new relation with the specified name and attributes. Relation names and attribute names should not be more than 20 characters long (including the NULL that terminates them), and you must enforce this. The format for each attribute is one of three — "i" for integer, "f" for floating point, or "sN" for character string of maximum length N (where N is a string of digits). The upper bound on N is 50. It is an error to create a relation that already exists.

The routine creates a Unix file for the relation and makes appropriate entries in the catalogs. It adds a record to the relation catalog describing the new relation, and adds records to the attribute catalog for all the attributes of the new relation.

7. *Destroy(argc, argv)*:

  argv[0]   = "destroy"
  argv[1]   = relation name
  argv[argc] = NIL

  This routine removes the specified relation from the database. Information about the relation is removed from the catalogs as well, and the associated Unix file is deleted.

8. *Load(argc, argv)*:

  argv[0]   = "load"
  argv[1]   = relation name
  argv[2]   = data file name
  argv[argc] = NIL

  This routine loads the specified relation, which must already have been created, with data from the specified file (the full path name of the file will be provided). Data files are in binary format as a sequence of data records — the difference between a data file and a relation is that the data file has nothing in it except the data itself, whereas a relation has additional page structure (for example, the slot maps). The sizes and types of the fields of the incoming records are ascertained by consulting the catalog information for the relation being loaded. It is an error to attempt to load data into a relation that is not empty. You may assume that the input file will not contain duplicate tuples.

9. *Print(argc, argv)*:

  argv[0]   = "print"
  argv[1]   = relation name
  argv[argc] = NIL

  This routine prints out the contents of the specified relation in the form of a table. The attribute values must be printed in a nice formatted way.

10. *BuildIndex(argc, argv)*:

  argv[0]   = "buildindex"
  argv[1]   = relation name
  argv[2]   = attribute name
  argv[argc] = NIL

  This routine creates a B+ tree index on the given attribute for the specified relation. To simplify the implementation, it is required that the desired indexes on a relation be created before any tuples are loaded. That is, the BuildIndex command should appear before the corresponding Load command.

11. *DropIndex(argc, argv)*:

  argv[0]   = "dropindex"
  argv[1]   = relation name
  argv[2]   = attribute name
  argv[argc] = NIL

  This routine destoys the index on the given attribute for the specified relation. If no attribute name is specified, the routine destroys *all* the indexes existing on the relation.

5   PHYSICAL LAYER

The Physical Layer consists of routines that manage files of records, since each relation in the database is stored as a file. This section describes the structure of files and records, and the meta-data catalogs.

5.1   Files and Records

For this project, all the records in each file have the same size. (Different files may have different record sizes, of course.) In MINIREL, files and records are implemented in the following manner: All files are structured as Unix files and each file consists of a number of pages depending on the number of records in the relation. Each page has several control fields plus a sequence of slots for storing records. The control fields include a *slot use map*, with one bit per record slot, indicating whether each slot is "in use" (has a record stored in it) or "free" (has no record currently stored in it). Record identifiers (RID's) are represented as 4-byte values — the upper 2 bytes are the page number and the lower 2 bytes are the record slot number within the page. A page identifier is represented as a RID with a nil (0) value in its record number portion, and nil page identifiers have 0's in both portions. The actual number of records per page depends on how the relation's record size compares to the size of a page. MINIREL uses $PAGESIZE = 512$ as the operating assumption. Records should not cross page boundaries.

5.2   Catalogs

A relational database has a number of catalogs that are used for storing meta-data information about relations, attributes, secondary indexes, views, user access privileges, etc. In MINIREL, a *relation catalog* is used for storing information about relations in the database, and an *attribute catalog* is used for storing information about attributes of each relation. These catalogs are themselves implemented as relations.

5.2.1   Relation Catalog

You should design the set of attributes for the relation catalog. A required minimal set of attributes is given below:

relName	—	Name of the relation described by this catalog record
recLength	—	Length of the relation's records, in bytes
recsPerPg	—	Number of records per page for the relation
numAttrs	—	Number of attributes for the relation
numRecs	—	Number of records currently in the relation
numPgs	—	Number of pages in the relation

Since the relation catalog is itself a relation, it should contain, as its first record, a description of itself! The relation catalog is called **relcat**. The second record in the relation catalog is for the attribute catalog, which is called **attrcat**.

5.2.2   Relation Cache

In addition to storing the above information, it is important for achieving reasonable performance to have this information quickly accessible for all currently open relations. For this reason, MINIREL maintains an array of relation catalog entries of current interest in main memory, that is, a *relation catalog cache*. No more than *MAXOPEN* = 20 relations are allowed to be open at the same time. Cache entries contain all of the fields from the relation catalog records, *plus* the following information:

relcatRid	—	RID for record holding this relation catalog record
relFile	—	File descriptor for the open relation
dirty	—	Set to true if the corresponding catalog record on disk becomes outdated
attrList	—	Linked list of attribute descriptors

When a relation is opened, its relation catalog record is copied into the cache for easy access. The *relcatRid* for the entry is saved to make it easy to update (if necessary) the relation catalog record on disk when the relation is subsequently closed. The *dirty* field is set to true if the cached version of the entry changes, and it is false until then. If this value is true when the relation is closed, the relation catalog record on disk is updated before the cache entry is reused. The *relFile* field is used to store the Unix file descriptor while the file is open. The role of the last field, *attrList*, is explained in Section 5.2.4.

After a relation is opened, MINIREL routines typically use the array index into the cache for the relation, called its **relation number**, for quick access of information about the relation. You are free to add, to the cache record structure, any additional information that will result in performance enhancements.

5.2.3   Attribute Catalog

The attribute catalog is analogous to the relation catalog — it contains a record for *every* attribute in *every* relation in the database, including both the relation catalog and itself. You should design the set of fields for the attribute catalog. A required minimal set of attributes is given below:

offset	—	offset of attribute within record
length	—	length of attribute
type	—	attribute type: "i", "f", or "s"
attrName	—	name of attribute
relName	—	name of relation

5.2.4 Attribute Cache

MINIREL makes the attribute information for open files readily available by cacheing this information in the same place that it caches their relation catalog information — in the relation cache. Since the number of attributes that relations have is not fixed, a linked list of attribute catalog entries is built. When a relation is opened, the attribute catalog is searched and the information for each of its attributes is read, adding each one in turn to the list in the relation's cache entry.

5.3 Buffer Management

MINIREL uses an array of memory buffers for storing pages on their way to and from the disk. There is *one buffer per open relation*, so the buffer pool is indexed by the *relation number* of the relation. The page processing routines operate using the buffers. When a page is requested, it is either found in the buffer associated with the relation being read, or else it is read in after replacing the page currently in the buffer (writing it to disk first if it's *dirty*). Finally, when a relation is closed, if it still has a *dirty* page in the buffer pool, that page is written to disk before the close is performed.

You are free to use a few additional buffers if needed in the implementation. However, these should be limited in number – in particular, you <u>cannot</u> assume that entire relations or entire indexes can be brought into memory.

5.4 Command Description

The physical layer involves writing a set of routines to manipulate the catalogs, files, pages and records. These routines are described in detail below.

5.4.1 Catalog Routines

1. *CreateCats()* — Create the system catalogs and place entries in the catalogs for the catalogs themselves.

2. *OpenCats()* — Open the system catalogs and place their entries into the catalog caches. Initialize the buffer pool control fields.

3. *CloseCats()* — Close the system catalogs (after closing any remaining open relations).

5.4.2 File Routines

1. *OpenRel(relName)* — Open relation *relName* and return its relation number.
   char *relName – Name of the relation to be opened.

2. *CloseRel(relNum)* — Close the relation identified by *relNum* (after writing to disk any changes in its buffer page or in its cached catalog information).
   int relNum – Relation number.

3. *FindRelNum(relName)* — Find the relation number of relation *relName* if it is currently open.
   char *relName – Name of the relation to be found in the cache.

5.4.3 Page Routines

1. *ReadPage(relNum, pid)* — Read page *pid* of the open relation specified by *relNum* and place the resulting *PAGESIZE* bytes in the open relation's buffer slot. Of course, if the page is already in the buffer, then don't actually read it.
   int relNum – Relation number.
   short pid – Page identifier.

2. *FlushPage(relNum)* — Flush the open relation's buffer pool page. This involves writing the page to disk if it's *dirty*, then marking the buffer page as being clean.
   int relNum – Relation number.

Note: Only the page routines should use I/O calls provided by C. All record routines should use the page routines as their I/O primitives.

5.4.4 Record Routines

1. *InsertRec(relNum, recPtr)* — Insert the record whose contents is pointed at by *recPtr* into relation *relNum*.
   Do not simply do a sequential scan of a file when you need to find a slot to insert a new record. Your solution
   must be more sophisticated, so that its performance will not be terrible.
   int relNum – Relation number.
   char *recPtr – A pointer to a record-sized byte array whose contents will be copied to an empty record slot in
   the relation.

2. *GetNextRec(relNum, startRid, foundRid, recPtr)* — Similar to *FindRec*, but without the search criteria speci-
   fication. Simply find the next record in the file (for sequential scan purposes).
   int relNum – Relation number.
   Rid *startRid – The record identifier from which to begin the search for the next record. The search will begin
   with the first record *after* startRid.
   Rid *foundRid – The record identifier of the first record after the one specified by startRid. If no record is
   found in the relation which follows startRid, foundRid will be returned with a null value.
   char *recPtr – A pointer to a record-sized byte array into which the contents of the next record (if any) will
   be put.

3. *FindRec(relNum, startRid, foundRid, recPtr, attrType, attrSize, attrOffset, valuePtr, compOp)* — Starting at
   record *startRid* in relation *relNum*, find the RID of the next record that meets the specification and put it in
   foundRid; if none is found, return *foundRid* = 0. Also, put the contents of the record found into the byte array
   pointed at by *recPtr*. The specification is that the value of the given attribute (of type *attrType*, size *attrSize*,
   and offset *attrOffset*) satisfies the relationship specified by the comparison operator, *compOp*, with the value
   pointed to by *valuePtr*. If *startRid* is 0, find the RID of first matching record in the file.
   int relNum – Relation number.
   Rid *startRid – The record identifier from which to begin the search for the record with the given attribute.
   The search will begin with the first record *after* startRid.
   Rid *foundRid – The record identifier of the first record after that specified by startRid which has the attribute
   described by attrType, attrSize, attrOffset, and valuePtr. If no matching record is found in the relation,
   foundRid will be returned with value null.
   char *recPtr – A pointer to a record-sized byte array into which the contents of the matching record (if any)
   will be put.
   char attrType – The data type (integer, float, or string) of the search attribute.
   int attrSize – The length of the search attribute, in bytes.
   int attrOffset – The offset into the record at which the search attribute is located.
   char *valuePtr – A pointer to a byte array which contains the search value.
   int compOp – Comparison operator.

4. *WriteRec(relNum, recPtr, recRid)* — Make the record-sized byte array pointed at by *recPtr* to be the new
   contents of record *recRid* in relation *relNum*.
   int relNum – Relation number.
   char *recPtr – A pointer to a record-sized byte array which contains the new contents of the record.
   Rid *recRid – The record identifier of the record which is to be updated.

5. *DeleteRec(relNum, recRid)* — Delete record *recRid* from relation *relNum*.
   int relNum – Relation number.
   Rid *rid – The record identifier of the record which is to be deleted.

5.5   Index Routines

MINIREL uses indexing to speed up execution of queries. The indexes are based on B^+ trees, with each index
represented by a Unix file and each tree node represented by a page in the file. You should implement key search,
insert and delete operations on the index, as well as index maintenance operations such as splitting and merging of
nodes (recursively up the tree if necessary).

   The structure of the index files and their manipulation routines are to be designed by you. The detailed description
given in this document of the design methodology for the other parts of MINIREL will help you in this process.

6   Error Handling

You should make your software as robust as possible. Check the return values on all routine calls and report any
errors via appropriate error messages if anything goes wrong. Basically, your code should be fairly bullet-proof, and

it should certainly never fail in a way that could compromise the data in the database!

All error messages are printed by **one** central error routine, the *ErrorMsgs* routine. This routine is implemented as part of the Physical Layer and has the following minimal interface:

ErrorMsgs(errorNum, printFlag)
int errorNum – Error message number.
int printFlag – Print out the error message if non-zero.

All errors should have numbers, and these numbers should be declared as constants in a definition module that all of your modules import. The *ErrorMsgs* routine is called when leaving routines in which an error has occurred — the call is of the form "return(*ErrorMsgs(...)*)". If *printFlag* is true, the routine prints an error message and then returns the value *errorNum*; otherwise it just returns *errorNum*.

The error messages must be *helpful*, that is, they should convey to the user not only what the problem is, but also suggest possible solutions for fixing the problem.

7 Implementation Notes

1. C aligns record fields on word boundaries. You need to take this alignment into account while computing sizes of structures.

2. Make full use of the typecasting feature of C.

3. Although users should not be allowed to directly modify the meta-data, they should be allowed to *view* this information through the logical layer routines.

4. Implement the indexing routines *after* you have completed the basic code in all three layers.

5. The names of the catalogs and their attributes should be *exactly* as specified in Section 5.2.

6. Ensure that there are no interactive parts in your code. Also, all error messages should go to *stdout*, not *stderr*.

7. For the prespecified minimal set of attributes of the catalogs, you should use exactly the same names (including case) as specified in Sections 5.2.1 and 5.2.3 .

8 Documentation

Your MINIREL software should be well-documented. Documentation includes program documentation, layer documentation, module documentation, function documentation, and in-line documentation.

For program documentation, write up a one page summary which describes the project, the input/output interface, the optimizations that have been incorporated, etc. Also list what you have NOT implemented (for example, indexing).

9 Course Files

All the project-related files are in `http://dsl.serc.iisc.ernet.in/~course/SW/MINIREL/minirel.tgz`.
The following files are available, with their locations as shown:

1. Project description files (MINIREL/doc)

2. Front-end files (MINIREL/frontend)

3. Basic include files (MINIREL/include)

4. Execution files (MINIREL/run)

5. Sample makefiles (MINIREL/template)

6. Program documentation example files (MINIREL/template)

7. Sample query files (MINIREL/query)

8. Binary data files (MINIREL/data/binary)

9. Ascii version of data files (MINIREL/data/ascii)

The frontend (parser) is MINIREL/frontend/FES.o . In order to use the parser, you need a main program. This is MINIREL/run/main.c . To create your final minirel executable, a sample makefile is available in MINIREL/run. If you invoke *make minirel* with this makefile, it links together the object files from each layer with the parser and the main program, and creates an executable called *minirel*. Your queries will be run on this executable. Note that MINIREL is a <u>single</u> Unix process.

Sample data files are in MINIREL/data/binary and sample query files are in MINIREL/query. To enable you to check that your *Load* routine is correctly reading in the binary data files, the ascii equivalents of these data files are provided in MINIREL/data/ascii.

10   Final Remarks

Read through this document <u>several</u> times and make sure you have a full understanding of the complete project *before* you start writing your code and designing your data structures. Discuss all design decisions with your partner to ensure that your code integrates cleanly. Split the project work with your partner *vertically*, not on a layer basis. You are encouraged to make additions and improvements to the implementation method suggested here (for example, by adding fields to the catalog cache), so long as you meet the basic interface specifications.

Good luck!